

# Statistical Machine Learning

Christoph Lampert



*Institute of Science and Technology*

Spring Semester 2013/2014 // Lecture 4

## (Generalized) Maximum Margin Classifiers – Optimization II

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^i$$

subject to

$$\begin{aligned} y^i (\langle w, x^i \rangle + b) &\geq 1 - \xi^i, \quad \text{for } i = 1, \dots, n, \\ \xi^i &\geq 0. \quad \text{for } i = 1, \dots, n. \end{aligned}$$

How to solve numerically?

- off-the-shelf Quadratic Program (QP) solver  
only for small dimensions and training sets (a few hundred),
- variants of gradient descent,  
high dimensional data, large training sets (millions)
- by convex duality,  
for very high dimensional data and not so many examples ( $d \gg n$ )

## Subgradient-Based Optimization

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^i$$

subject to

$$y^i (\langle w, x^i \rangle + b) \geq 1 - \xi^i, \quad \text{and} \quad \xi^i \geq 0, \quad \text{for } i = 1, \dots, n.$$

## Subgradient-Based Optimization

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^i$$

subject to

$$y^i(\langle w, x^i \rangle + b) \geq 1 - \xi^i, \quad \text{and} \quad \xi^i \geq 0, \quad \text{for } i = 1, \dots, n.$$

For any fixed  $(w, b)$  we can find the optimal  $\xi_1, \dots, \xi_n$ :

$$\xi_i = \max\{ 0, 1 - y_i(\langle w, x_i \rangle + b) \}.$$

## Subgradient-Based Optimization

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^i$$

subject to

$$y^i(\langle w, x^i \rangle + b) \geq 1 - \xi^i, \quad \text{and} \quad \xi^i \geq 0, \quad \text{for } i = 1, \dots, n.$$

For any fixed  $(w, b)$  we can find the optimal  $\xi_1, \dots, \xi_n$ :

$$\xi_i = \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}.$$

Plug into original problem:

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}.$$

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}.$$

- unconstrained optimization problem
- convex
  - ▶  $\frac{1}{2} \|w\|^2$  is convex (differentiable with Hessian = Id  $\succcurlyeq$  0)
  - ▶ linear/affine functions are convex
  - ▶ pointwise **max** over convex functions is convex.
  - ▶ sum of convex functions is convex.
- *not differentiable!*

# SVM Training in the Primal

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}.$$


- unconstrained optimization problem
- convex
  - ▶  $\frac{1}{2} \|w\|^2$  is convex (differentiable with Hessian = Id  $\succcurlyeq$  0)
  - ▶ linear/affine functions are convex
  - ▶ pointwise **max** over convex functions is convex.
  - ▶ sum of convex functions is convex.
- *not differentiable!*

We can't use gradient descent, since some points have no gradients!

# Subgradients

**Definition:** Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be a **convex** function. A vector  $v \in \mathbb{R}^d$  is called a **subgradient** of  $f$  at  $w_0$ , if

$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$






# Subgradients

**Definition:** Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be a **convex** function. A vector  $v \in \mathbb{R}^d$  is called a **subgradient** of  $f$  at  $w_0$ , if


$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$



# Subgradients

**Definition:** Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be a **convex** function. A vector  $v \in \mathbb{R}^d$  is called a **subgradient** of  $f$  at  $w_0$ , if


$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$



# Subgradients

**Definition:** Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be a **convex** function. A vector  $v \in \mathbb{R}^d$  is called a **subgradient** of  $f$  at  $w_0$ , if


$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$



# Subgradients

**Definition:** Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be a **convex** function. A vector  $v \in \mathbb{R}^d$  is called a **subgradient** of  $f$  at  $w_0$ , if

$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$




A general convex  $f$  can have be different subgradients at a position.

- We write  $\nabla f(w_0)$  for the set of subgradients of  $f$  at  $w_0$ ,
- $v \in \nabla f(w_0)$  indicates that  $v$  is a subgradient of  $f$  at  $w_0$ .

# Subgradients

- For differentiable  $f$ , the gradient  $v = \nabla f(w_0)$  is the only subgradient.



- If  $f_1, \dots, f_K$  are differentiable at  $w_0$  and


$$f(w) = \mathbf{max}\{f_1(w), \dots, f_K(w)\},$$

then  $v = \nabla f_k(w_0)$  is a subgradient of  $f$  at  $w_0$ , where  $k$  any index for which  $f_k(w_0) = f(w_0)$ .

- Subgradients are only well defined for convex functions!


## Illustration: Optimization using Gradients

$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$




## Illustration: Optimization using Gradients

$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$



## Illustration: Optimization using Gradients


$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$






## Illustration: Optimization using Gradients

$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$




## Illustration: Optimization using Gradients

$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$




## Illustration: Optimization using Gradients

$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$



## Illustration: Optimization using Gradients


$$f(w_1, w_2) = (w_1)^2 + 2(w_2)^2 \quad \text{convex, differentiable}$$



Gradient of a differentiable function is a **descent direction**:


- for any  $w_t$  there exists an  $\eta$  such that  $f(w_t + \eta v) < f(w_t)$

## Illustration: Optimization using Subgradients?




## Illustration: Optimization using Subgradients?

$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$




## Illustration: Optimization using Subgradients?

$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$



## Illustration: Optimization using Subgradients?


$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$






## Illustration: Optimization using Subgradients?

$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$




## Illustration: Optimization using Subgradients?

$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$



## Illustration: Optimization using Subgradients?

$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$




Subgradient might not be a **not a descent direction**:

- for  $w_t$  we might have  $f(w_t + \eta v) \geq f(w_t)$  for all  $\eta \in \mathbb{R}$

## Illustration: Optimization using Subgradients?

$$f(w_1, w_2) = |w_1| + 2|w_2| \quad \text{convex, not differentiable}$$



Subgradient might not be a **not a descent direction**:

- for  $w_t$  we might have  $f(w_t + \eta v) \geq f(w_t)$  for all  $\eta \in \mathbb{R}$
- but: there is an  $\eta$  that brings us closer to the optimum,  
 $\|w_{t+1} - w^*\| < \|w_t - w^*\|$  (Proof: exercise...)

## Subgradient Method (not Descent!)

**input** step sizes  $\eta_1, \eta_2, \dots$

1:  $w_1 \leftarrow 0$

2: **for**  $t = 1, \dots, T$  **do**

3:    $v \leftarrow$  a subgradient of  $\mathcal{L}$  at  $w_t$

4:    $w_{t+1} \leftarrow w_t - \eta_t v$

5: **end for**

**output**  $w_t$  with smallest values  $\mathcal{L}(w_t)$  for  $t = 1, \dots, T$

## Subgradient Method (not Descent!)

**input** step sizes  $\eta_1, \eta_2, \dots$

1:  $w_1 \leftarrow 0$

2: **for**  $t = 1, \dots, T$  **do**

3:    $v \leftarrow$  a subgradient of  $\mathcal{L}$  at  $w_t$

4:    $w_{t+1} \leftarrow w_t - \eta_t v$

5: **end for**

**output**  $w_t$  with smallest values  $\mathcal{L}(w_t)$  for  $t = 1, \dots, T$

Stepsize rules: how to choose  $\eta_1, \eta_2, \dots, ?$

- $\eta_t = \eta$  constant: will get us (only) close to the optimum
- decrease slowly, but not too slowly: converges to optimum

$$\sum_{t=1}^{\infty} \eta_t = \infty \quad \sum_{t=1}^{\infty} (\eta_t)^2 < \infty \quad \text{e.g. } \eta_t = \frac{\eta}{t + t_0}$$

How to choose overall  $\eta$ ? trial-and-error

- Try difference values, see which one decreases the objective (fastest)

Many objective functions in ML contain a sum over all training examples:

$$\mathcal{L}_{LogReg}(w) = \sum_{i=1}^n \log(1 + \exp(-y_i(\langle w, x_i \rangle + b))),$$

$$\mathcal{L}_{SVM}(w) = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}.$$

Computing the gradient or subgradient scales like  $O(nd)$ ,

- $d$  is the dimensionality of the data
- $n$  is the number of training examples.

Both  $d$  and  $n$  can be big (millions). What can we do?

- we'll not get rid of  $O(d)$ , since  $w \in \mathbb{R}^d$ ,
- but we can get rid of the scaling with  $O(n)$  for each update!

Let  $f(w) = \sum_{i=1}^n f_i(w)$ , with convex, differentiable  $f_1, \dots, f_n$ .

## Stochastic Gradient Descent

**input** step sizes  $\eta_1, \eta_2, \dots$

1:  $w_1 \leftarrow 0$

2: **for**  $t = 1, \dots, T$  **do**

3:    $i \leftarrow$  random index in  $1, 2, \dots, n$

4:    $v \leftarrow n \nabla f_i(w_t)$

5:    $w_{t+1} \leftarrow w_t - \eta_t v$

6: **end for**

**output**  $w_T$ , or average  $\frac{1}{T-T_0} \sum_{t=T_0}^T w_t$

- Each iteration takes only  $O(d)$ ,
- No line search, since evaluating  $f(w - \eta v)$  would be  $O(nd)$ ,
- Objective does not decrease in every step,
- Converges to optimum if  $\eta_t$  is square summable, but not summable.



Let  $f(w) = \sum_{i=1}^n f_i(w)$ , with differentiable  $f_1, \dots, f_n$ .

## Stochastic Subgradient Method

**input** step sizes  $\eta_1, \eta_2, \dots$

1:  $w_1 \leftarrow 0$

2: **for**  $t = 1, \dots, T$  **do**

3:    $i \leftarrow$  random index in  $1, 2, \dots, n$

4:    $v \leftarrow$   $n$  times a subgradient of  $f_i$  at  $w_t$

5:    $w_{t+1} \leftarrow w_t - \eta_t v$

6: **end for**

**output**  $w_T$ , or average  $\frac{1}{T-T_0} \sum_{t=T_0}^T w_t$

- Each iteration takes only  $O(d)$ ,
- Converges to optimum if  $\eta_t$  is square summable, but not summable.
- Often better not to pick completely at random, but go in epochs: randomly shuffle dataset, go through all examples, reshuffle, etc.

## Stochastic Primal SVMs Training

$$\mathcal{L}_{SVM}(w) = \sum_{i=1}^n \left( \frac{1}{2n} \|w\|^2 + C \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\} \right).$$

**input** step sizes  $\eta_1, \eta_2, \dots$  or step size rule, such as  $\eta_t = \frac{\eta}{t+t_0}$

1:  $(w_1, b_1) \leftarrow (0, 0)$

2: **for**  $t = 1, \dots, T$  **do**

3:   pick  $(x, y)$  from  $\mathcal{D}$  (randomly, or in epochs)

4:   **if**  $y\langle x, w \rangle + b \geq 1$  **then**

5:      $w_{t+1} \leftarrow (1 - \eta_t)w_t$

6:   **else**

7:      $w_{t+1} \leftarrow (1 - \eta_t)w_t + nC\eta_t yx$

8:      $b_{t+1} \leftarrow \eta_t nC y$

9:   **end if**

10: **end for**

**output**  $w_T$ , or average  $\frac{1}{T-T_0} \sum_{t=T_0}^T w_t$

State-of-the-art in SVM training, but setting stepsizes can be painful.

Back to the original formulation

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^i$$

subject to, for  $i = 1, \dots, n$ ,

$$y^i(\langle w, x^i \rangle + b) \geq 1 - \xi^i, \quad \text{and} \quad \xi^i \geq 0.$$

Convex optimization problem: we can study its **dual problem**.

## General Principle of Dualization

Assume a constrained optimization problem:

$$\min_{\theta \in \Theta} f(\theta)$$

subject to

$$g_1(\theta) \leq 0, \quad g_2(\theta) \leq 0, \quad \dots, \quad g_k(\theta) \leq 0.$$

## General Principle of Dualization

Assume a constrained optimization problem:

$$\min_{\theta \in \Theta} f(\theta)$$

subject to

$$g_1(\theta) \leq 0, \quad g_2(\theta) \leq 0, \quad \dots, \quad g_k(\theta) \leq 0.$$

We define the **Lagrangian**, that combines objective and constraints

$$\mathcal{L}(\theta, \alpha) = f(\theta) + \alpha_1 g_1(\theta) + \dots + \alpha_k g_k(\theta)$$

with **Lagrange multipliers**,  $\alpha_1, \dots, \alpha_k$ , such that

$$\max_{\alpha_1 \geq 0, \dots, \alpha_k \geq 0} \mathcal{L}(\theta, \alpha) = \begin{cases} f(\theta) & \text{if } g_1(\theta) \leq 0, \quad g_2(\theta) \leq 0, \quad \dots, \quad g_k(\theta) \leq 0 \\ \infty & \text{otherwise.} \end{cases}$$

Any optimal solution,  $\theta$ , for  $\min_{\theta \in \Theta} \max_{\alpha \geq 0} \mathcal{L}(\theta, \alpha)$  is also optimal for the original constrained problem.

## Theorem (Special Case of Slater's Condition)

*If  $f$  is convex,  $g_1, \dots, g_k$  are affine functions, and there exists at least one point  $\theta \in \text{relint}(\Theta)$  that is feasible (i.e.  $g_i(\theta) \leq 0$  for  $i = 1, \dots, k$ ). Then*

$$\min_{\theta \in \Theta} \max_{\alpha \geq 0} \mathcal{L}(\theta, \alpha) = \max_{\alpha \geq 0} \min_{\theta \in \Theta} \mathcal{L}(\theta, \alpha)$$

# General Principle of Dualization

## Theorem (Special Case of Slater's Condition)

If  $f$  is convex,  $g_1, \dots, g_k$  are affine functions, and there exists at least one point  $\theta \in \text{relint}(\Theta)$  that is feasible (i.e.  $g_i(\theta) \leq 0$  for  $i = 1, \dots, k$ ). Then

$$\min_{\theta \in \Theta} \max_{\alpha \geq 0} \mathcal{L}(\theta, \alpha) = \max_{\alpha \geq 0} \min_{\theta \in \Theta} \mathcal{L}(\theta, \alpha)$$

We call  $f(\theta)$  the **primal function** and  $h(\alpha) = \min_{\theta \in \Theta} \mathcal{L}(\theta, \alpha)$  be the **dual function**.

The theorem states that minimizing the primal  $f(\theta)$  (with constraints given by the  $g_k$ ) is equivalent to maximizing its dual  $h(\alpha)$  (with  $\alpha \geq 0$ ).

## Dualizing of the SVM optimization problem

The SVM optimization problem fulfills the conditions of the theorem.

$$\min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^i$$

subject to, for  $i = 1, \dots, n$ ,

$$y^i(\langle w, x^i \rangle + b) \geq 1 - \xi^i, \quad \text{and} \quad \xi^i \geq 0.$$

We can compute its minimal value as  $\max_{\alpha \geq 0, \beta \geq 0} h(\alpha, \beta)$  with

$$h(\alpha, \beta) = \min_{(w, b)} \quad \frac{1}{2} \|w\|^2 + C \sum_i \xi_i + \sum_i \alpha_i (1 - \xi_i - y^i(\langle w, x^i \rangle + b)) - \sum_i \beta_i \xi_i$$

(Blackboard...)



## Dualizing of the SVM optimization problem

In a minimum w.r.t.  $(w, b)$ :

$$0 = \frac{\partial}{\partial w} \mathcal{L}(w, b, \xi, \alpha, \beta) = w - \sum_i \alpha_i y^i x^i \quad \Rightarrow \quad w = \sum_i \alpha_i y^i x^i$$

$$0 = \frac{\partial}{\partial b} \mathcal{L}(w, b, \xi, \alpha, \beta) = \sum_i \alpha_i y^i$$

$$0 = \frac{\partial}{\partial \xi_i} \mathcal{L}(w, b, \xi, \alpha, \beta) = C - \alpha_i - \beta_i$$

Insert new constraints into objective:

$$\max_{\alpha \geq 0} \frac{1}{2} \left\| \sum_i \alpha_i y^i x^i \right\|^2 + \sum_i \alpha_i - \sum_i \alpha_i y_i \left\langle \sum_j \alpha_j y^j x^j, x^i \right\rangle$$

## SVM Dual Optimization Problem

$$\max_{\alpha \geq 0} \quad -\frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y^i y^j \langle x^i, x^j \rangle + \sum_i \alpha_i$$

subject to  $\sum_i \alpha_i y_i = 0$  and  $0 \leq \alpha_i \leq C$ , for  $i = 1, \dots, n$ .

- Examples  $x^i$  with  $\alpha_i \neq 0$  are called **support vectors**.
- From the coefficients  $\alpha_1, \dots, \alpha_n$  we can recover the optimal  $w$ :

$$w = \sum_i \alpha_i y^i x^i$$

$$b = 1 - y^i \langle x^i, w \rangle \quad \text{for any } i \text{ with } 0 < \alpha_i < C$$

(more complex rule for  $b$  if not such  $i$  exists).

- The prediction rule becomes

$$g(x) = \text{sign} \left( \langle w, x \rangle + b \right) = \text{sign} \left( \sum_{i=1}^n \alpha_i y_i \langle x_i, x \rangle + b \right)$$

# SVM Dual Optimization Problem

$$\max_{\alpha \geq 0} \quad -\frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y^i y^j \langle x^i, x^j \rangle + \sum_i \alpha_i$$

subject to

$$\sum_i \alpha_i y_i = 0 \quad \text{and} \quad 0 \leq \alpha_i \leq C, \quad \text{for } i = 1, \dots, n.$$

Why solve the dual optimization problem?

- fewer unknowns:  $\alpha \in \mathbb{R}^n$  instead of  $(w, b, \xi) \in \mathbb{R}^{d+1+n}$
- less storage when  $d \gg n$ :  
 $(\langle x^i, x^j \rangle)_{i,j} \in \mathbb{R}^{n \times n}$  instead of  $(x^1, \dots, x^n) \in \mathbb{R}^{n \times d}$
- **Kernelization**

## Definition (Positive Definite Kernel Function)

Let  $\mathcal{X}$  be a non-empty set. A function  $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is called **positive definite kernel function**, if the following conditions hold:

- $k$  is symmetric, i.e.  $k(x, x') = k(x', x)$  for all  $x, x' \in \mathcal{X}$ .
- For any finite set of points  $x_1, \dots, x_n \in \mathcal{X}$ , the *kernel matrix*

$$K_{ij} = (k(x_i, x_j))_{i,j} \quad (1)$$

is positive semidefinite, i.e. for all vectors  $t \in \mathbb{R}^n$

$$\sum_{i,j=1}^n t_i K_{ij} t_j \geq 0. \quad (2)$$

## Lemma (Kernel function)

Let  $\phi : \mathcal{X} \rightarrow \mathcal{H}$  be a feature map into a Hilbert space  $\mathcal{H}$ . Then the function

$$k(x, \bar{x}) = \langle \phi(x), \phi(\bar{x}) \rangle_{\mathcal{H}}$$

is a positive definite kernel function.

### Proof.

- symmetry:  $k(x, \bar{x}) = \langle \phi(x), \phi(\bar{x}) \rangle_{\mathcal{H}} = \langle \phi(\bar{x}), \phi(x) \rangle_{\mathcal{H}} = k(\bar{x}, x)$
- positive definiteness:  $x_1, \dots, x_n \in \mathcal{X}$ , and arbitrary  $t \in \mathbb{R}^n$ , then

$$\begin{aligned} \sum_{i,j=1}^n t_i k(x_i, x_j) t_j &= \sum_{i,j=1}^n t_i t_j \langle \phi(x^i), \phi(x^j) \rangle_{\mathcal{H}} \\ &= \left\langle \sum_i t_i \phi(x^i), \sum_j t_j \phi(x^j) \right\rangle_{\mathcal{H}} = \left\| \sum_i t_i \phi(x^i) \right\|_{\mathcal{H}}^2 \geq 0. \end{aligned}$$



## Theorem (Mercer's Condition)

*Let  $\mathcal{X}$  be non-empty set. For any positive definite kernel function  $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ , there exists a Hilbert space  $\mathcal{H}$  with inner product  $\langle \cdot, \cdot \rangle_{\mathcal{H}}$ , and a feature map  $\phi : \mathcal{X} \rightarrow \mathcal{H}$  such that*

$$k(x, \bar{x}) = \langle \phi(x), \phi(\bar{x}) \rangle_{\mathcal{H}}.$$

**Proof.** later, in more refined form

Note:  $\mathcal{H}$  and  $\phi$  are not unique, e.g.

$$k(x, \bar{x}) = 2x\bar{x}$$

- $\mathcal{H}_1 = \mathbb{R}$ ,  $\phi_1(x) = \sqrt{2}x$ ,  $\langle \phi_1(x), \phi_1(\bar{x}) \rangle_{\mathcal{H}_1} = 2x\bar{x}$
- $\mathcal{H}_2 = \mathbb{R}^2$ ,  $\phi_2(x) = \begin{pmatrix} x \\ -x \end{pmatrix}$ ,  $\langle \phi_2(x), \phi_2(\bar{x}) \rangle_{\mathcal{H}_2} = 2x\bar{x}$
- $\mathcal{H}_3 = \mathbb{R}^3$ ,  $\phi_3(x) = \begin{pmatrix} x \\ 0 \\ x \end{pmatrix}$ ,  $\langle \phi_3(x), \phi_3(\bar{x}) \rangle_{\mathcal{H}_3} = 2x\bar{x}$ , etc.

## Definition (Reproducing Kernel Hilbert Space)

Let  $\mathcal{H}$  be a Hilbert space of functions  $f : \mathcal{X} \rightarrow \mathbb{R}$ . A kernel  $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is called **reproducing kernel**, if

$$f(x) = \langle k(x, \cdot), f(\cdot) \rangle_{\mathcal{H}} \quad \text{for all } f \in \mathcal{H}.$$

$\mathcal{H}$  is then called a **reproducing kernel Hilbert space (RKHS)**.

## Theorem (Moore-Aronszajn Theorem)

*Let  $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  be a positive definite kernel on  $\mathcal{X}$ . Then there is a unique Hilbert space of functions,  $f : \mathcal{X} \rightarrow \mathbb{R}$ , for which  $k$  is a reproducing kernel.*

**Proof sketch.** One can construct the space explicitly: Set

$$\mathcal{H}^{\text{pre}} = \text{span}\{ k(\cdot, x) \text{ for } x \in \mathcal{X} \},$$

i.e., for every  $f \in \mathcal{H}^{\text{pre}}$  exist  $x^1, \dots, x^m \in \mathcal{X}$  and  $\alpha^1, \dots, \alpha^m \in \mathbb{R}$ , with

$$f(\cdot) = \sum_{i=1}^m \alpha^i k(\cdot, x^i).$$

We define an inner product as

$$\langle f, g \rangle = \left\langle \sum_i \alpha^i k(\cdot, x^i), \sum_j \bar{\alpha}^j k(\cdot, \bar{x}^j) \right\rangle := \sum_{i,j} \alpha^i \bar{\alpha}^j k(x^i, \bar{x}^j).$$

Make  $\mathcal{H}^{\text{pre}}$  into Hilbert space  $\mathcal{H}$  by enforcing *completeness*.

Complete proof: [B. Schölkopf, A. Smola, "*Learning with Kernels*", 2001].



Let

- $\mathcal{D} = \{(x^1, y^1), \dots, (x^n, y^n)\} \subset \mathcal{X} \times \{\pm 1\}$  training set
- $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  be a pos.def. kernel with feature map  $\phi : \mathcal{X} \rightarrow \mathcal{H}$ .

## Support Vector Machine in Kernelized Form

For any  $C > 0$ , the max-margin classifier for the feature map  $\phi$  is

$$g(x) = \text{sign } f(x) \quad \text{with} \quad f(x) = \sum_i \alpha_i k(x^i, x) + b,$$

for coefficients  $\alpha_1, \dots, \alpha_n$  obtained by solving

$$\min_{\alpha^1, \dots, \alpha^n \in \mathbb{R}} \quad -\frac{1}{2} \sum_{i,j=1}^n \alpha^i \alpha^j y^i y^j k(x^i, x^j) + \sum_{i=1}^n \alpha^i$$

subject to  $\sum_i \alpha_i y_i = 0$  and  $0 \leq \alpha_i \leq C$ , for  $i = 1, \dots, n$ .

Note: we don't need to know  $\phi$  or  $\mathcal{H}$ , explicitly. Knowing  $k$  is enough.

## Useful and Popular Kernel Functions

For  $x, \bar{x} \in \mathbb{R}^d$ :

- $k(x, \bar{x}) = (1 + \langle x, x' \rangle)^p$  for  $p \in \mathbb{N}$  (*polynomial kernel*)  
 $f(x) = \sum_i \alpha_i y^i k(x^i, x)$  = polynomial of degree  $d$
- $k(x, \bar{x}) = \exp(-\lambda \|x - \bar{x}\|^2)$  for  $\lambda > 0$  (*Gaussian or RBF kernel*)  
 $f(x) = \sum_i \alpha_i y^i \exp(-\lambda \|x^i - x\|^2)$  = weighted/soft nearest neighbor

For  $x, \bar{x}$  histograms with  $d$  bins:

- $k(x, \bar{x}) = \sum_{j=1}^d \mathbf{min}(x_j, \bar{x}_j)$  histogram intersection kernel
- $k(x, \bar{x}) = \sum_{j=1}^d \frac{x_j \bar{x}_j}{x_j + \bar{x}_j}$   $\chi^2$  kernel
- $k(x, \bar{x}) = \exp \left( -\lambda \sum_{j=1}^d \frac{(x_j - \bar{x}_j)^2}{x_j + \bar{x}_j} \right)$  exponentiated  $\chi^2$  kernel

Generally: interpret kernel function as a **similarly measure**.